Quadratic Programming

Julia Wrobel

Overview

Today, we cover:

- Finish linear programming
- Quadratic programming
- Examples of LP and QP in statistics

Announcements

- Final project due

Strong duality

- We know that at feasible solutions for both dual and primal, the solution to the dual is always greater or equal
- If there is a difference between the largest primal value and the smallest dual value it is called the "**Duality gap**".
    - When are these values equal?

Strong Duality Theorem: If the primal has an optimal solution, then the dual also has an optimal solution and there is no duality gap, i.e.:

$$\text{objective of dual} \;=\; \text{objective of primal.}$$

- Condition for strong duality: the feasible region must be nonempty and bounded.
- The **duality gap** is zero when strong duality holds.

Complementary slackness theorem

The **complementary slackness theorem** provides a necessary condition for optimality in linear programming. It states that:

If x^* and y^* are feasible solutions of primal and dual problems, then x^* and y^* are both optimal if and only if

1. $y^{*T}(b - Ax^*) = 0$
2. $(y^{*T}A - c)x^* = 0$

This theorem is the foundation of the **interior point method** class of LP solvers.

Complementary slackness theorem

Consider a simple LP problem:

$$\begin{aligned} \textit{max} \quad & z = x_1 + x_2, \\ \textit{s.t.} \quad & x_1 + 2x_2 \leq 100, \\ & 2x_1 + x_2 \leq 100, \\ & x_1, x_2 \geq 0. \end{aligned}$$

Its dual problem is:

Complementary slackness theorem

Consider a simple LP problem:

$$\begin{aligned} \textit{max} \quad & z = x_1 + x_2, \\ \textit{s.t.} \quad & x_1 + 2x_2 \leq 100, \\ & 2x_1 + x_2 \leq 100, \\ & x_1, x_2 \geq 0. \end{aligned}$$

Its dual problem is:

$$\begin{aligned} \textit{min} \quad & z = 100y_1 + 100y_2, \\ \textit{s.t.} \quad & y_1 + 2y_2 \geq 1, \\ & 2y_1 + y_2 \geq 1, \\ & y_1, y_2 \geq 0. \end{aligned}$$

Complementary slackness theorem

The complementary slackness theorem states:

Complementary slackness theorem

The complementary slackness theorem states:

$$
\begin{aligned}
y_1(100 - x_1 - 2x_2) &= 0, \\
y_2(100 - 2x_1 - x_2) &= 0, \\
x_1(y_1 + 2y_2 - 1) &= 0, \\
x_2(2y_1 + y_2 - 1) &= 0.
\end{aligned}
$$

At the optimal point, we have $\mathbf{x} = (\mathbf{100/3}, \mathbf{100/3})$, $\mathbf{y} = (\mathbf{1/3}, \mathbf{1/3})$. The complementary slackness holds.

- All the constraints are bounded in both primal and dual problems.

Plot of constraints

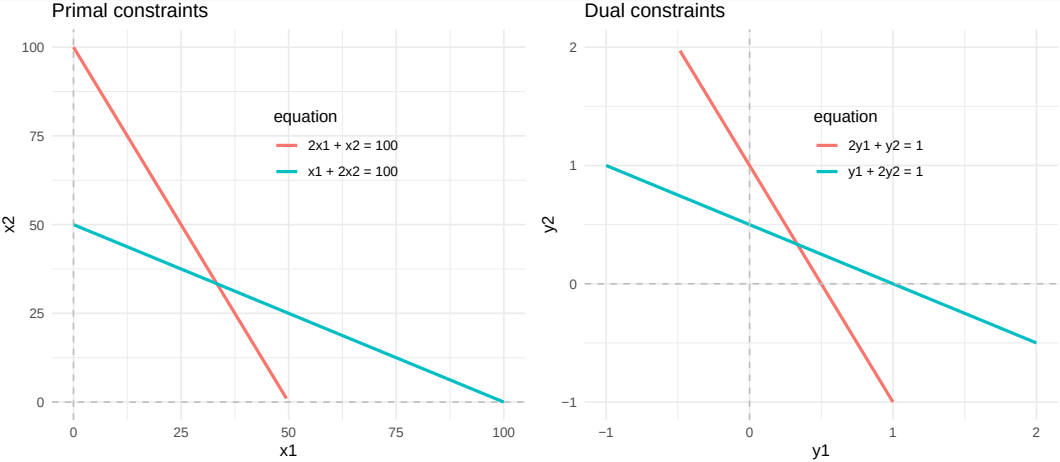

Complementary slackness theorem

Modify the simple LP problem a bit:

$$\begin{aligned} \textit{max} \quad & z = 3x_1 + x_2, \\ \textit{s.t.} \quad & x_1 + 2x_2 \leq 100, \\ & 2x_1 + x_2 \leq 100, \\ & x_1, x_2 \geq 0. \end{aligned}$$

and its dual problem is:

$$\begin{aligned} \textit{min} \quad & z = 100y_1 + 100y_2, \\ \textit{s.t.} \quad & y_1 + 2y_2 \geq 3, \\ & 2y_1 + y_2 \geq 1, \\ & y_1, y_2 \geq 0. \end{aligned}$$

Complementary slackness theorem

Now the complementary slackness theorem states:

$$
\begin{aligned}
y_1(100 - x_1 - 2x_2) &= 0, \\
y_2(100 - 2x_1 - x_2) &= 0, \\
x_1(y_1 + 2y_2 - 3) &= 0, \\
x_2(2y_1 + y_2 - 1) &= 0.
\end{aligned}
$$

For this LP problem, the optimal solutions are $\mathbf{x} = (\mathbf{50}, \mathbf{0})$, $\mathbf{y} = (\mathbf{0}, \mathbf{1.5})$. Complementary slackness still holds. We observe that:

- In the primal problem:
  - The first constraint is **inactive** (slack > 0) at the optimum, so its shadow price $y_1 = 0$
  - Second constraint is **active** (binding, slack $= 0$), so y_2 can be non-zero

Complementary slackness theorem

For this LP problem, the optimal solutions are $\mathbf{x} = (\mathbf{50}, \mathbf{0})$, $\mathbf{y} = (\mathbf{0}, \mathbf{1.5})$. Complementary slackness still holds. We observe that:

- In the primal problem:
  - The first constraint is **inactive** (slack > 0), so its shadow price $y_1 = 0$.
  - The second constraint is **active** (binding, slack $= 0$), so y_2 can be non-zero.
- In the dual problem:
  - The first constraint is **active** (binding), so its primal variable x_1 can be non-zero.
  - The second constraint is **inactive** (slack > 0), so its primal variable $x_2 = 0$.

Plot of constraints

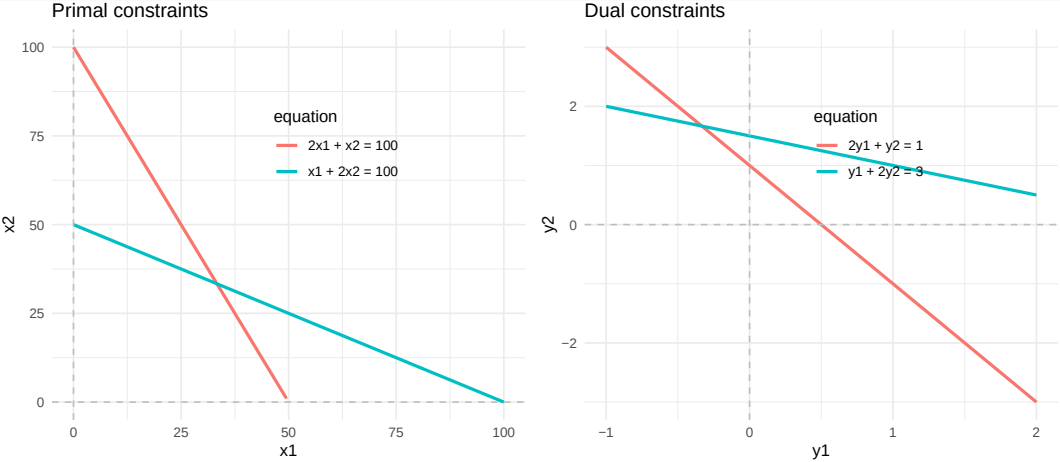

Economics interpretation

Dual variables can be called "**shadow prices**" of the primal constraint, how objective function would increase if constraint was relaxed.

- If a primal constraint is bounded, relaxing that constraint would result in a gain (improve the objective function), shadow price is non-zero.
- If a primal constraint is unbounded, relaxing that constraint would not improve the objective function, shadow price is zero.

Economics interpretation

Bakery example: Imagine you're running a bakery, and can make a limited number of cakes because you have a limited amount of flour. Shadow price tells you how much more profit you'd gain if you had a little more flour.

- If you're tight on flour, getting more flour would let you bake more cakes and make more money. - In this case, the shadow price is nonzero because the extra resource improves your outcome.
- If you have plenty of flour, adding more doesn't help—you're already making as many cakes as you can sell. In this case, the shadow price is zero because extra flour doesn't change your profit.

So, the dual variable (shadow price) tells you how valuable it would be to relax a constraint—meaning allowing more of that constrained resource. If relaxing the constraint helps, the shadow price is positive. If it doesn't make a difference, the shadow price is zero.

Primal/dual optimality conditions

Given the primal and dual problem with slack/surplus variables added:

Primal:

$$\begin{aligned} \textit{max} \quad & cx \\ \textit{s.t.} \quad & Ax + w = b, \quad x, w \geq 0. \end{aligned}$$

Dual:

$$\begin{aligned} \textit{min} \quad & b^T y \\ \textit{s.t.} \quad & A^T y - z = c^T, \quad y, z \geq 0. \end{aligned}$$

The *complementary slackness theorem* states that at optimal solution, we should have $x_j z_j = 0 \forall j$, and $w_i y_i = 0 \forall i$. In matrix notation, the complementary conditions can be rewritten as

$$XZe = 0, WYe = 0$$

- X, Z, W, Y are diagonal matrices
- e is a vector of 1's

Primal/dual optimality conditions

We now have the optimality conditions for the primal/dual problems as:

$$
\begin{aligned}
Ax + w - b &= 0, \\
A^T y - z - c^T &= 0, \\
XZe &= 0, \\
WYe &= 0, \\
x, y, w, z &\geq 0.
\end{aligned}
$$

- The first two conditions are simply the constraints for primal/dual problems
- The next two are complementary slackness
- The last one is the non-negativity constraint

Primal/dual optimality conditions

We now have the optimality conditions for the primal/dual problems as:

$$
\begin{aligned}
Ax + w - b &= 0, \\
A^T y - z - c^T &= 0, \\
XZe &= 0, \\
WYe &= 0, \\
x, y, w, z &\geq 0.
\end{aligned}
$$

Ignoring the non-negativity constraints, this is a set of $2n + 2m$ equations with $2n + 2m$ unknowns (n and m are the number of unknowns and constraints in the primal problem), which can be solved using Newton's method.

Primal-dual interior point method

The primal-dual interior point method finds the primal-dual optimal solution (x^*, y^*, w^*, z^*) by applying Newton's method to the primal-dual optimality conditions.

- The direction and length of the steps are modified in each step so that the non-negativity condition is strictly satisfied in each iteration.

To be specific, define the following function $F : R^{2n+2m} \to R^{2n+2m}$:

$$\mathbf{F}(x, y, w, z) = \begin{bmatrix} \mathbf{A}x + \mathbf{w} - \mathbf{b} \\ \mathbf{A}^T y - \mathbf{z} - \mathbf{c}^T \\ XZe \\ WYe \end{bmatrix}$$

The goal is to find solution for $F = 0$.

Primal-dual interior point method

Applying Newton's method, if at iteration k the variables are (x^k, y^k, w^k, z^k), we obtain a search direction $(\delta x, \delta y, \delta w, \delta z)$ by solving the linear equations:

$$F'(x^k, y^k, w^k, z^k) \begin{bmatrix} \delta x \\ \delta y \\ \delta w \\ \delta z \end{bmatrix} = -F(x^k, y^k, w^k, z^k).$$

Primal-dual interior point method

Then the update will be:

$$(x^{k+1}, y^{k+1}, w^{k+1}, z^{k+1}) = (x^k, y^k, w^k, z^k) + \alpha(\delta x, \delta y, \delta w, \delta z)$$

with $\alpha \in (0, 1]$ chosen so that the result from the next iteration is feasible.

- Solving for $(\delta x, \delta y, \delta w, \delta z)$ using the equation on the previous slide enforces feasibility of solutions

An improved algorithm

The algorithm in its current setup is not ideal because often the steps are small to avoid violating the positivity constraints. Instead:

- The value of $XZe + WYe$ represents the duality gap
- Instead of trying to eliminate the duality gap (set it to 0), reduce the duality gap by some factor in each step

We replace the complementary slackness by:

$$XZe = \mu_x e \tag{1}$$
$$WYe = \mu_y e \tag{2}$$

When $\mu_x, \mu_y \to 0$ as $k \to \infty$ the solution from this system will converge to the optimal solution of the original LP problem. Good selections of μ are $\mu_x^{k+1} = (x^k)^T z^k / n$ and $\mu_y^{k+1} = (w^k)^T y^k / m$

- n and m are the dimensions of x and y

An improved algorithm

Under the new algorithm, at the kth iteration, the Newton equations become:

$$\begin{bmatrix} A & 0 & I & -0 \\ 0 & A^T & 0 & -I \\ Z & 0 & 0 & X \\ 0 & W & Y & 0 \end{bmatrix} \begin{bmatrix} \delta x \\ \delta y \\ \delta w \\ \delta z \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ -X^k Z^k e + \mu_x^k e \\ -W^k Y^k e + \mu_y^k e \end{bmatrix}$$

This provides the **general primal-dual interior point method** as follows:

1. Choose strictly feasible initial solution (x^0, y^0, w^0, z^0), and set $k = 0$. Then repeat (2) and (3) until convergence:
2. Solve system (1) to obtain the updates $(\delta x, \delta y, \delta w, \delta z)$
3. Update the solution $(x^{k+1}, y^{k+1}, w^{k+1}, z^{k+1}) = (x^k, y^k, w^k, z^k) + \alpha^k(\delta x, \delta y, \delta w, \delta z)$. α^k is chosen so that all variables are ≥ 0.

Lagrange multipliers

The method of Lagrange multiplier is a general algorithm for optimization problems with equality constraints. For example, consider a problem:

$$\begin{aligned} \max \quad & f(x, y) \\ s.t. \quad & g(x, y) = c \end{aligned}$$

We introduce a new variable λ called a **Lagrange multiplier** and form the following new objective function:

$$L(x, y, \lambda) = f(x, y) + \lambda[g(x, y) - c]$$

We will then optimize L with respect to x, y and λ using typical methods. Note that the condition $\partial L / \partial \lambda = 0$ at the optimal solution guarantees that the constraints will be satisfied.

The barrier problem

The interior point algorithm is closely related to the **Barrier Problem**. Go back to the primal problem:

$$\begin{aligned} \max \quad & z = cx \\ s.t. \quad & Ax + w = b; x, w \geq 0 \end{aligned}$$

The non-negativity constraints can be replaced by adding two **barrier terms** in the objective function. The barrier term is defined as $B(x) = \sum_j \log x_j$, which is finite as long as x_j is positive. Then the primal problem becomes:

$$\begin{aligned} \max \quad & z = cx + \mu_x B(x) + \mu_y B(w) \\ s.t. \quad & Ax + w = b \end{aligned}$$

The barrier terms make sure x and w won't become negative.

The barrier problem

The Lagrangian for this problem is (using y as the multiplier):

$$L(x, y, w) = cx + \mu_x B(x) + \mu_y B(w) + y^T(b - w - Ax)$$

The optimal solution for the problem satisfies:

$$c + \mu_x X^{-1} e - A^T y = 0 \tag{3}$$
$$b - w - Ax = 0 \tag{4}$$
$$\mu_y W^{-1} e - y = 0 \tag{5}$$

Define new variables $z = \mu_x X^{-1} e$ and rewrite these conditions, we obtain exactly the same set of equations as the relaxed optimality conditions for primal-dual problem. **Note**: $z = \mu_x X^{-1} e$ is also a condition so there are 4 in total.

Introduction to quadratic programming

We have discussed linear programming, where both the objective function and constraints are linear functions of the unknowns.

The **quadratic programming** (QP) problem has a quadratic objective function and linear constraints:

$$\begin{aligned} \max \quad & f(x) = \frac{1}{2}x^T Bx + cx \\ s.t. \quad & Ax \leq b, x \geq 0 \end{aligned}$$

The algorithm for solving QP problem is very similar to that for LP. But first we need to first introduce the **KKT conditions**.

KKT conditions

For LP, complementary slackness gave us conditions for optimality. **KKT conditions** generalize this to non-linear objective functions. They are necessary conditions for a solution to be optimal in a general non-linear programming problem.

Consider the following problem:

$$\begin{aligned} \max \quad & f(x) \\ s.t. \quad & g_i(x) \leq 0, i = 1, \dots, I \\ & h_j(x) = 0, j = 1, \dots, J \end{aligned}$$

The Lagrangian is: $L(x, y, z) = f(x) - \sum_i y_i g_i(x) - \sum_j z_j h_j(x)$. Then at the optimal solution, the following **KKT** conditions must be satisfied:

- Primal feasibility: $g_i(x^*) \leq 0, h_j(x^*) = 0$
- Dual feasibility: $y_i \geq 0$
- Complementary slackness: $y_i g_i(x^*) = 0$
- Stationary: $\nabla f(x^*) - \sum_i y_i \nabla g_i(x) - \sum_j z_j \nabla h_j(x) = 0$

Optimal solution for QP

The Lagrangian for the QP problem is

$$L(x, \mu, \lambda) = \frac{1}{2}x^T Bx + cx - y^T(Ax - b) + z^T x$$

The KKT conditions for the QP problem are:

- Primal feasibility: $Ax \leq b, x \geq 0$
- Dual feasibility: $y \geq 0, z \geq 0$ (z multiplies the constraint $x \geq 0$, written as $-x \leq 0$, so it requires the same non-negativity as y)
- Complementary slackness: $Y(Ax - b) = 0, Zx = 0$
- Stationary: $Bx + c - A^T y + z = 0$

Here, Y and Z are diagonal matrices with y and z at the diagonal. This can be solved using an interior point method.

Optimal solution for QP

To be specific, define slack variable $w = b - Ax$, the optimality conditions become

$$
\begin{aligned}
Ax - b + w &= 0, \\
Bx + c - A^T y + z &= 0, \\
Zx &= 0, \\
Yw &= 0, \\
x, y, z, w &\geq 0.
\end{aligned}
$$

The unknowns are x, y, z, w. We can then obtain the Jacobians, form the Newton equation and solve for the optimal solution iteratively.

QP in R

The quadprog package provides functions (`solve.QP.compact`) to solve quadratic programming problem. Pay attention to the definition of function parameters. They are slightly different from what I have used in the standard form!

- Solve $\min \frac{1}{2}(x_1^2 + x_2^2)$ s.t. $2x_1 + x_2 \geq 1$:

```
library(quadprog)
Dmat = diag(rep(1,2))
dvec = rep(0,2)
Amat = matrix(c(2,1))
solve.QP(Dmat=Dmat, dvec=dvec, Amat=Amat, bvec=c(1))
```

Review

We have now covered:

- LP problem set up
- Simplex methods
- Duality
- Interior point algorithms
- Quadratic programming

Now you should be able to formulate an LP/QP problem and solve it. But how are these useful in statistics?

- Remember that LP is an optimization algorithm
- There are plenty of optimization problems in statistics, e.g. MLE
- It's just a matter of formulating the objective function and constraints.

Lab problem

Write out the primal and dual for the carpenter problem.

- Check the set of feasible solutions (corner points).
  - What are the values of the objective functions of the primal and dual at each vertex?
  - Does weak duality appear to hold?

Quantile regression

The goal of regression is to understand the relationship between an outcome and covariates. Traditional regression models the conditional mean, $E(Y|X)$

- Relies on normal error assumptions and homoscedasticity
- Sensitive to outliers and skewed distributions

Quantile regression:

- Provides a more exhaustive description of the data
- The collection of regressions at all quantiles would give a complete picture of outcome-covariate relationships

Quantile regression

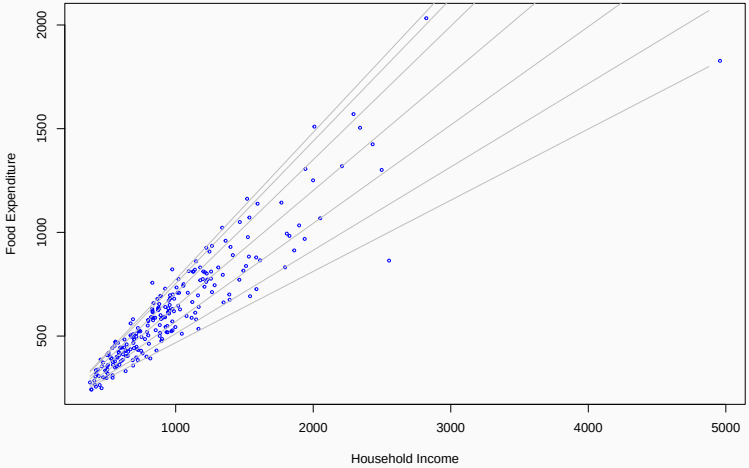

Colorado cannabis and driving study

Colorado cannabis and driving study

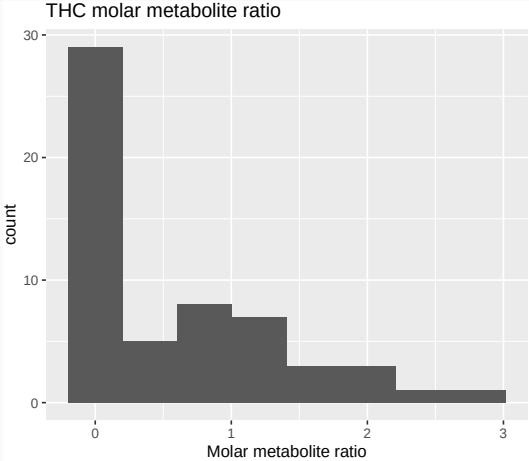

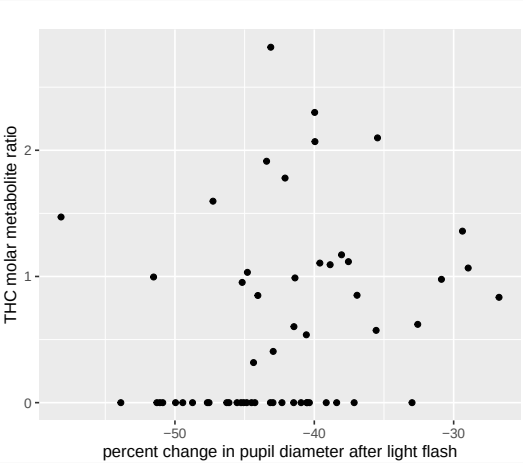

Quantile regression model

Regress conditional quantiles of response on the covariates. Assume the outcome Y is continuous.

- Classical model: $Q_\tau(Y|X) = X\beta_\tau$
- $Q_\tau(Y|X)$ is the τth conditional quantile of Y given X

The above model is equivalent to specifying

$$Y = X\beta_\tau + \epsilon; \quad Q_\tau(\epsilon|X) = 0.$$

In comparison, mean regression is:

$$Y = X\beta + \epsilon; \quad E(\epsilon|X) = 0.$$

Pros and cons

Advantages:

- Regression at a sequence of quantiles provides a more complete view of data
- Inference is robust to outliers
- Estimation is more efficient when residual normality is highly violated
- Allows interpretation in the outcome's original scale of measurement

Disadvantages:

- To be useful, needs to regress on a set of quantiles: computational burden
- Solution has no closed form
- Adaptation to non-continuous outcomes is difficult

The loss function

Link between estimands and loss functions:

- To obtain sample mean of $\{y_1, y_2, \dots y_n\}$ minimize $\sum_i (y_i - b)^2$
- To obtain sample median of $\{y_1, y_2, \dots y_n\}$ minimize $\sum_i |y_i - b|$

It can be shown that to obtain the sample τth quantile, one needs to minimize asymmetric absolute loss, that is, compute

$$\hat{Q}_\tau = \arg\min_b \left\{ \sum_{i:y_i \geq b} \tau |y_i - b| + \sum_{i:y_i < b} (1-\tau)|y_i - b| \right\}.$$

For notational simplicity, define $\rho_\tau(x) = x[\tau - (x < 0)]$, where $(x < 0)$ is an indicator equal to 1 if $x < 0$ and 0 otherwise. This compactly encodes the asymmetric weighting: $\rho_\tau(x) = \tau x$ when $x \geq 0$ and $\rho_\tau(x) = (\tau - 1)x$ when $x < 0$.

The loss function

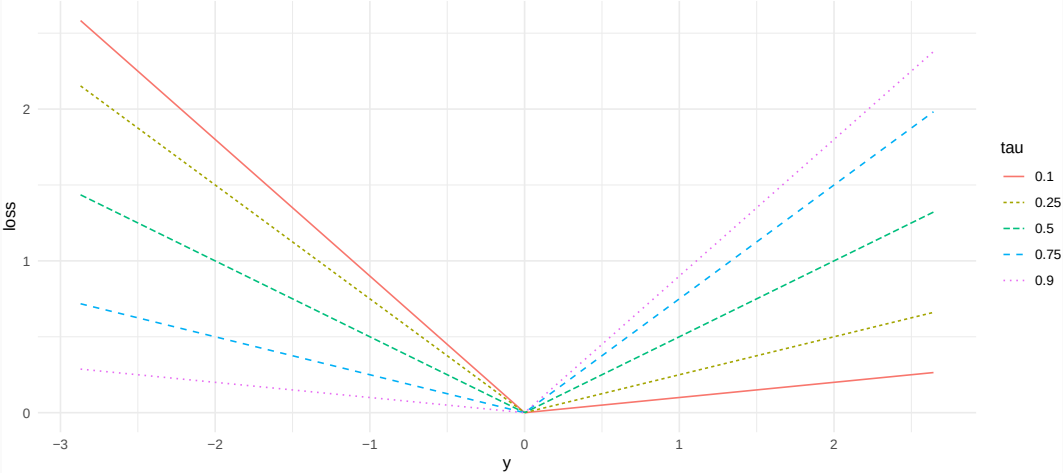

Estimator

The linear quantile regression model is fitted by determining

$$\hat{\beta}_\tau = \arg\min_b \sum_{i=1}^n \rho_\tau(y_i - x_i b)$$

The estimator has all "expected" properties:

- Scale equivariance:

$$\hat{\beta}_\tau(ay, X) = a\hat{\beta}_\tau(y, X), \quad \hat{\beta}_\tau(-ay, X) = -a\hat{\beta}_{1-\tau}(y, X)$$

- Shift (or regression) equivariance: $\hat{\beta}_\tau(y + \gamma, X) = \hat{\beta}_\tau(y, X) + \gamma$
- Equivariance to reparametrization of design: $\hat{\beta}_\tau(y, XA) = A^{-1}\hat{\beta}_\tau(y, X)$

Asymmetric Double Exponential (ADE) distribution

Density function for ADE is

$$f(y;\mu,\sigma,\tau) = \frac{\tau(1-\tau)}{\sigma}\exp\left\{-\rho_\tau(\frac{y-\mu}{\sigma})\right\}$$

- Least squares estimator $\iff$ MLE if residuals are Normal
- QR estimator $\iff$ MLE if residuals are ADE
- If residuals re iid $ADE(0,1,\tau)$, then the log-likelihood for β_τ is

$$l(\beta_\tau;Y,X,\tau) = -\sum_i^n \rho_\tau(y_i - x_i\beta_\tau) + c_0$$

QR model fitting

The QR estimation problem $\hat{\beta}_\tau = \arg\min_b \sum_{i=1}^n \rho_\tau(y_i - x_i b)$ can be framed as an LP problem!

- First, define a set of new variables

$$u_i \equiv [y_i - x_i b]_+ \tag{6}$$
$$v_i \equiv [y_i - x_i b]_- \tag{7}$$
$$b_+ \equiv [b]_+ \tag{8}$$
$$b_- \equiv [b]_- \tag{9}$$

Note:

- $[\cdot]_+$ and $[\cdot]_-$ mean the positive and negative part of a number
- $x \equiv [x]_+ - [x]_-$
  - $[x]_+ = \max(x, 0)$
  - $[x]_- = \max(-x, 0)$

QR model fitting

Remember the objective function is $\sum_{i=1}^{n} \rho_\tau(y_i - x_i b)$. Notice that

- When $y_i - x_i b \geq 0$, we have $\rho_\tau(y_i - x_i b) = \tau u_i$, and $v_i = 0$.
- When $y_i - x_i b < 0$, we have $\rho_\tau(y_i - x_i b) = (1-\tau) v_i$, and $u_i = 0$.

So we can write $\rho_\tau(y_i - x_i b) = \tau u_i + (1-\tau) v_i$.

QR model fitting

The optimization problem can then be reformulated as:

$$
\begin{aligned}
\max \quad & -\sum_{i=1}^{n}[\tau u_i + (1-\tau)v_i] \\
\text{s.t.} \quad & y_i = x_i b_+ - x_i b_- + u_i - v_i \\
& u_i, v_i \geq 0, \quad i = 1, \dots, n \\
& b_+, b_- \geq 0
\end{aligned}
$$

This is a standard LP problem can be solved by simplex or interior point methods.

QR model fitting, cont

Written in matrix notation, and make u_i, v_i, b_+, b_- as unknowns, get

$$\begin{aligned} \max \quad & -[\mathbf{0}, \mathbf{0}, \tau, \mathbf{1}-\tau] \begin{bmatrix} \mathbf{b}_+ \\ \mathbf{b}_- \\ \mathbf{u} \\ \mathbf{v} \end{bmatrix} \\ s.t. \quad & [\mathbf{X}, -\mathbf{X}, \mathbf{I}, -\mathbf{I}] \begin{bmatrix} \mathbf{b}_+ \\ \mathbf{b}_- \\ \mathbf{u} \\ \mathbf{v} \end{bmatrix} = \mathbf{y} \\ & \mathbf{b}_+, \mathbf{b}_-, \mathbf{u}, \mathbf{v} \geq 0 \end{aligned}$$

The dual problem is:

$$\begin{aligned} \min \quad & \mathbf{y}^T \mathbf{d} \\ s.t. \quad & \begin{bmatrix} \mathbf{X}^T \\ -\mathbf{X}^T \\ \mathbf{I} \\ -\mathbf{I} \end{bmatrix} \mathbf{d} \geq - \begin{bmatrix} \mathbf{0} \\ \mathbf{0} \\ \tau \\ \mathbf{1}-\tau \end{bmatrix} \end{aligned}$$

QR model fitting, cont

Manipulating the dual constraints, get

$$\begin{aligned} \mathbf{X}^T\mathbf{d} &= \mathbf{0} \\ -\tau \le \mathbf{d} &\le 1 - \tau \end{aligned}$$

Define a new variable $\mathbf{a} = 1 - \tau - \mathbf{d}$, the original LP problem can be formulated as:

$$\begin{aligned} \max \quad & \mathbf{y}^T\mathbf{a} \\ s.t. \quad & \mathbf{X}^T\mathbf{a} = \mathbf{X}^T(1-\tau) \\ & 0 \le \mathbf{a} \le 1 \end{aligned}$$

Adding slack variables $\mathbf{s}$ for the $\le$ constraints, the problem can be formulated in the standard form, and can be solved by using either Simplex or Interior Point methods.

$$\begin{aligned} \max \quad & \mathbf{y}^T\mathbf{a} \\ s.t. \quad & \mathbf{X}^T\mathbf{a} = \mathbf{X}^T(1-\tau) \\ & \mathbf{a} + \mathbf{s} = 1 \\ & \mathbf{a}, \mathbf{s} \ge 0 \end{aligned}$$

QR model fitting, cont

In this form,

- $\mathbf{y}$ (n-vector) are outcomes
- $\mathbf{X}$ ($n \times p$ matrix) are predictors,
- $\mathbf{a}$ and $\mathbf{s}$ (n-vectors) are unknowns.
- $1 - \tau$ and 1 need to be n-vectors

There are $2n$ unknowns and $p + n$ constraints.

Once we have optimal $\mathbf{a}$, $\mathbf{d}$ can be obtained given τ (the quantile). Then depending on which constraints are hit in the dual problem, one can determine the set of basic variables in primal problem, and then solve for the primal variables (just solve a set of linear equations), and then obtain β.

QR in R

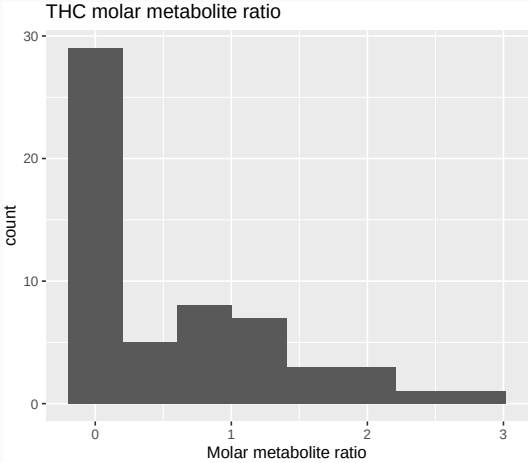
THC molar metabolite ratio
count
30
20
10
0
0
1
2
3
Molar metabolite ratio

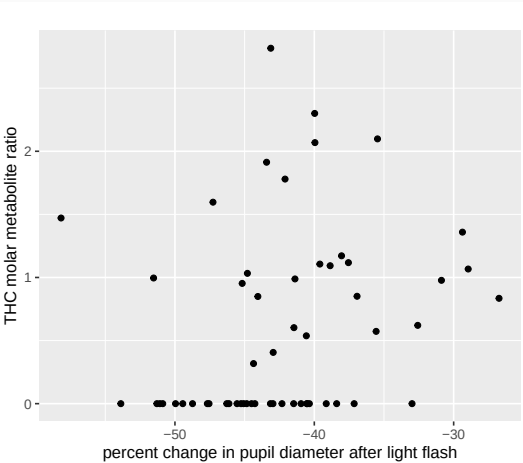
THC molar metabolite ratio
2
1
0
−50
−40
−30
percent change in pupil diameter after light flash

QR in R

Want to compare linear and quantile regression

```r
ols = lm(t_mmr1 ~ p_change, data = cannabis)

# default method, which is a simplex method
median_regression = rq(t_mmr1 ~ p_change, tau = 0.5, data = cannabis,
                       method = "br")
qr.9 = rq(t_mmr1 ~ p_change, tau = 0.9, data = cannabis,
          method = "br")

summary(ols)$coefficients
```

```
              Estimate Std. Error  t value    Pr(>|t|)
(Intercept) 1.80873656 0.65925047 2.743626 0.008188134
p_change    0.02884236 0.01540728 1.871995 0.066528189
```

```r
summary(median_regression, se = "boot")$coefficients
```

QR in R

Blue is ols, red is median regression, pink is tau = 0.9

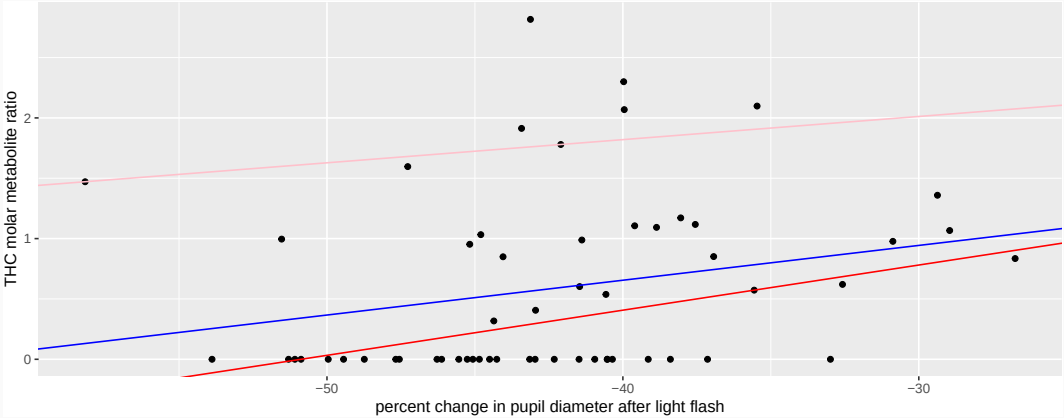

LASSO

Consider usual regression settings with data $(\mathbf{x_i}, \mathbf{y_i})$, where $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})$ is a p-vector of predictors and y_i is the response for the ith subject. The OLS setting finds a coefficient vector to minimize the residual sum of squares:

$$\hat{\beta} = \arg\min_{\beta} \sum_{i}^{n} (y_i - x_i\beta)^2$$

- Solution is the MLE assuming a Normal model:

$$y_i = \mathbf{x}_i\beta + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2)$$

- This is undesirable when p is large because

1. When p>n, model can't be fit
2. One typically wants a more parsimonius model for interpretability

LASSO

LASSO stands for "Least Absolute Shrinkage and Selection Operator", which aims for model selection when p is large (works even $p > n$). The LASSO procedure will "shrink" the coefficients toward 0 and eventually force some to be exactly zero

- predictors with $\beta = 0$ will be selected out

The LASSO estimates are defined as:

$$\tilde{\beta} = \arg\min_{\beta} \left\{ \sum_{i}^{n} (y_i - \mathbf{x}_i \beta)^2 \right\} s.t. ||\beta||_1 \leq \lambda$$

- $||\beta||_1 = \sum_{j=1}^{p} |b_j|$ is the L_1 norm
- $\lambda \geq 0$ is a tuning parameter controlling the strength of shrinkage

What kind of optimization problem does this look like?

Regularization side note

LASSO is a regularization approach- regularization prevents overfitting by adding a penalty term to the objective function.

- LASSO tries to minimize the residual sum of square, with a constraint on the sum of the absolute values of the coefficients.
- There are other types of regularized regressions
- **Ridge regression**: adds an L_2 penalty, $\sum_j b_j^2 \leq \lambda$

LASSO vs. Ridge visualized

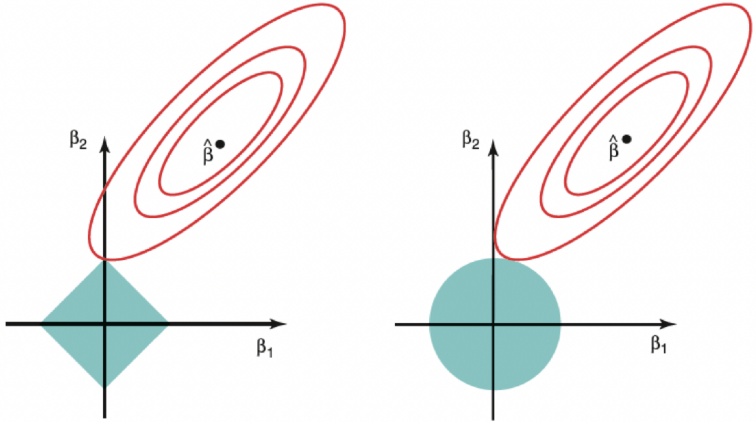

LASSO model fitting

The LASSO problem can be solved by quadratic programming algorithms.

$$\max \quad -\sum_{i}^{n}\left(y_i - \sum_j \beta_j x_{ij}\right)^2 \tag{10}$$

$$s.t. \quad \sum_j |b_j| \leq \lambda \tag{11}$$

- **Issue**: not in standard LP/QP form, since the constraints have the absolute value operator

LASSO model fitting

The trick is to convert the problem into the standard QP problem setting, i.e., to remove the absolute value operator. This is the **same decomposition used in QR** model fitting (u_i, v_i for residuals), now applied to the coefficients β_j.

- Let $\beta_j = \beta_j^+ - \beta_j^-$, where $\beta_j^+, \beta_j^- \geq 0$ (positive and negative parts)
- Then $|\beta_j| = \beta_j^+ + \beta_j^-$ and the problem can be written as:

$$
\begin{aligned}
\max \quad & -\sum_i^n \left(y_i - \sum_j \beta_j^+ x_{ij} + \sum_j \beta_j^- x_{ij} \right)^2 \\
s.t. \quad & \sum_j (\beta_j^+ + \beta_j^-) \leq \lambda, \\
& \beta_j^+, \beta_j^- \geq 0
\end{aligned}
$$

This is now a standard QP problem can be solved by standard QP solvers.

A little more on LASSO

The Lagrangian for the LASSO optimization problem is:

$$L(,\lambda) = -\sum_{i=1}^{n}\left(y_i - \sum_j \beta_j x_{ij}\right)^2 - \lambda\sum_{j=1}^{p}|\beta_j|$$

This is equivalent to the likelihood function for a Bayesian model with a double exponential (DE) prior on βs (remember ADE used in quantile regression?):

$$Y|X,\beta \sim N(X\beta,\sigma^2)$$
$$\beta_j \sim DE(1/\lambda)$$

where the DE density function is

$$f(x,\tau) = \frac{1}{2\tau}\exp\left(\frac{-|x|}{\tau}\right)$$

LASSO in R

The package `glmnet` is standard for elastic net and LASSO in R.

```
### Lasso results from R package "glmnet"
library(glmnet)
lasso_mod = glmnet(x = dplyr::select(cannabis, -id, -t_mmr1),
                   y = cannabis$t_mmr1)
plot(lasso_mod, "lambda")
```

LASSO in R

The package `glmnet` is standard for elastic net and LASSO in R.

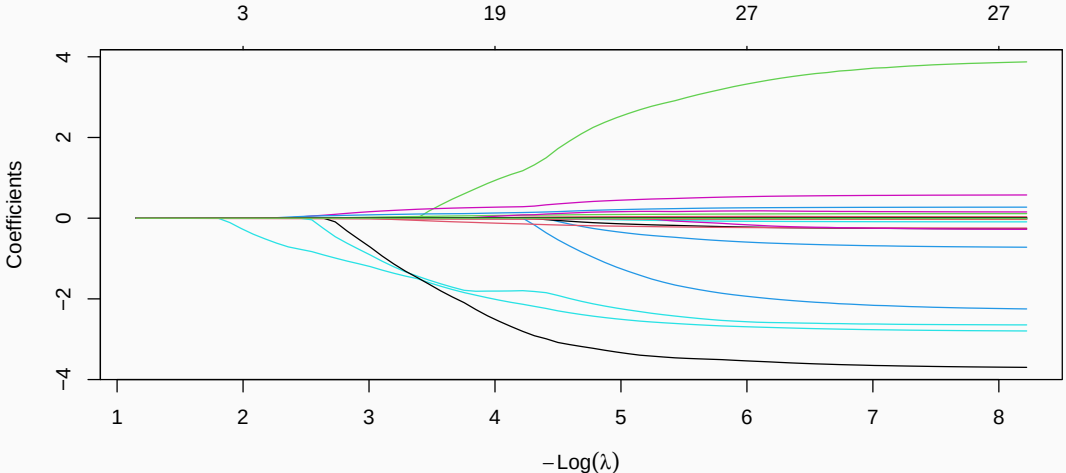

Resources

- Hao Wu's lecture with matrix represenation of Simplex algorithm

Some videos

- convexity explained, and duality
- KKT conditions explained
- Simple explanation of Lagrangian